

Data & ML Platform Engineer
Databricks
San Francisco, CA

About the Role:

We are looking for a Data & ML Platform Engineer to build scalable infrastructure powering data intelligence, distributed machine learning, and analytics at Databricks.

Responsibilities:

- Architect resilient data platforms using Apache Spark, Delta Lake, and Ray.
- Build and scale managed MLflow and Kubeflow platforms for distributed training.
- Design cloud infrastructure using Docker, Kubernetes, and AWS.
- Write clean, performant services in Python, Scala, and Java.

Requirements:

- Minimum 3 years of software engineering experience.
- Proficiency in Python, Java, Scala, and SQL.
- Hands-on experience with Apache Spark, Delta Lake, or Ray.
- Experience managing workloads with Docker, Kubernetes, and AWS.
- Strong familiarity with Git and modern CI/CD practices.

Preferred Qualifications:

- Production experience with MLflow, Kubeflow, or distributed model training.
- Deep understanding of cloud data architectures and distributed storage engines.