

Title: Research Engineer, Foundation Models
Google DeepMind
London, UK

Role Overview:

Google DeepMind is looking for a Research Engineer to join our Foundation Models team to advance frontier artificial intelligence and large-scale deep learning research.

Responsibilities:

- Implement, train, and evaluate frontier foundation models and generative architectures.
- Scale deep learning training pipelines across high-performance TPU and GPU accelerators.
- Collaborate with research scientists to prototype novel algorithms and model architectures.

Minimum Qualifications:

- 2+ years of experience in research engineering or machine learning.
- Strong programming background in Python.
- Extensive experience with deep learning frameworks including PyTorch, JAX, or TensorFlow.
- Practical understanding of deep learning fundamentals, transformers, and neural network optimization.

Preferred Qualifications:

- Experience with large-scale multi-node distributed model training.
- Publications or contributions to open-source foundation model repositories.